CSCI 347 - Project 3 Report

Preston Duffield
Western Washington University
May 29, 2023

Edge Detector

Edge Detector is an edge detection program that uses a laplacian filter to highlight the edges in a set of images. The images are specified as command line arguments and the project is designed to process multiple images concurrently using threads. The goal of the project is to concurrently apply the laplacian filter to multiple images, improving the execution time compared to sequential execution.

The program contains several functions, structs, and a global variable, all of which work together to process the images:

1. `PPMPixel` Struct: This struct represents a pixel in an image, with the RGB values of the pixel.

2. `parameter` Struct: This struct holds the information each thread needs to perform its task.

3. `file_name_args` Struct: This struct is used to hold the input and output filenames for each image.

4. `total_elapsed_time` Global Variable: This is a double value that tracks the total time taken by all threads to process all input images.

5. `compute_laplacian_threadfn`: This is the thread function. It applies the Laplacian filter to a section of the image.

6. `apply_filters`: This function uses threads to apply the Laplacian filter to an image and measures the time taken.

7. `write_image`: This function saves a processed image to a file.

8. `read_image`: This function reads an image file and returns a PPMPixel pointer to the pixel data.

9. `manage_image_file`: This is another thread function. It manages the entire process of reading an image file, processing the image, and writing the result to a file.

10. `main`: This is the entry point of the program. It checks command line arguments, creates threads to process each image file, and prints the total elapsed time.

Shell Script

Edge detector can be run from a shell script. The script, `run_program.sh`, accepts a directory as a command line argument. It then collects all of the `.ppm` files and passes them as input into the edge detection program.

Concurrency

Concurrency in this program is achieved through the use of POSIX threads (pthreads). For each image file passed as a command line argument, a new thread is created in the main function. Each of these threads runs the `manage_image_file` function, which manages the processing of one image.

Within the `manage_image_file` function, the `apply_filters` function is called, which further divides the work among a specified number of threads (defined by the `LAPLACIAN_THREADS` macro). Each of these threads runs the `compute_laplacian_threadfn` function, which applies the Laplacian filter to a section of the image.

In terms of handling possible race conditions, the main area where this could occur is when updating the `total_elapsed_time` global variable. However, as each image file is processed by a separate thread and the time taken to process each image is added to `total_elapsed_time` in a single operation, there's no risk of race conditions here as the operation is atomic. Furthermore, as the processed images are written to separate files, there is no risk of a race condition when writing the output.

Experiment

The following experiment was ran on a 2.2 GHz 6-Core Intel Core i7 processor. A Python program[1] was written to run `run_program.sh` 10 times and record the results to a .csv file. Each time the program was ran on 40 images of various sizes. 10 replicants (runs) were performed for each test and the times where noted. Using the same photos, run amounts, and processor, we can isolate our variable, which is the amount of threads that `edge_detector.c` opens. The experiement was ran with `LAPLACIAN_THREADS`, the value which defines the number of threads that work on a single image, set to values 1 through 16. The experiment was run 16 times, with 10 replicants each, for a total of 160 data points.

Experiemental Design Summary

Variables: The value of `LAPLACIAN_THREADS`, tested from 1 to 16.

Response: The combined time in seconds needed for each thread to complete its task.

Replicants: 10 replicants for each level of `LAPLACIAN_THREADS`, for a total of 160 data points.

Constants:

- 40 images were processed in each run.

- each test was run on the same processor.

---

[1]See Appendix

Results

Statistics

Variable	N	N*	Mean	SE Mean	StDev	Minimum	Q1	Median	Q3	Maximum
1	10	0	22.386	0.394	1.245	20.112	21.468	22.582	23.528	23.723
2	10	0	21.296	0.424	1.340	18.214	20.864	21.388	22.019	22.975
3	10	0	21.730	0.379	1.198	19.201	21.264	21.692	22.649	23.353
4	10	0	20.324	0.423	1.336	17.763	19.601	20.112	21.815	22.183
5	10	0	19.089	0.601	1.900	15.926	18.054	18.576	20.730	22.394
6	10	0	18.167	0.172	0.544	16.932	17.934	18.200	18.633	18.773
7	10	0	18.207	0.222	0.701	16.705	17.970	18.304	18.620	19.084
8	10	0	16.925	0.246	0.777	15.884	16.098	16.972	17.648	17.959
9	10	0	16.835	0.272	0.859	15.345	16.469	16.753	17.142	18.770
10	10	0	17.069	0.402	1.271	14.488	16.439	17.193	17.676	19.364
11	10	0	16.881	0.323	1.021	14.950	15.970	17.120	17.488	18.352
12	10	0	17.613	0.487	1.541	15.431	16.576	17.506	18.346	20.553
13	10	0	17.823	0.490	1.550	15.007	16.964	17.567	19.137	20.095
14	10	0	17.579	0.509	1.609	15.561	16.363	17.396	18.389	21.265
15	10	0	15.223	0.548	1.733	12.571	13.520	15.671	16.632	17.211
16	10	0	16.351	0.632	2.000	13.657	14.528	16.253	18.537	19.075

Figure 1: Descriptive statistics of the data.

The descriptive statistics of the data show that the means are grouped around the values of around 16 to 23 seconds. Note that "N" is the amount of replicants, and "Variable" is the amount of laplacian threads.
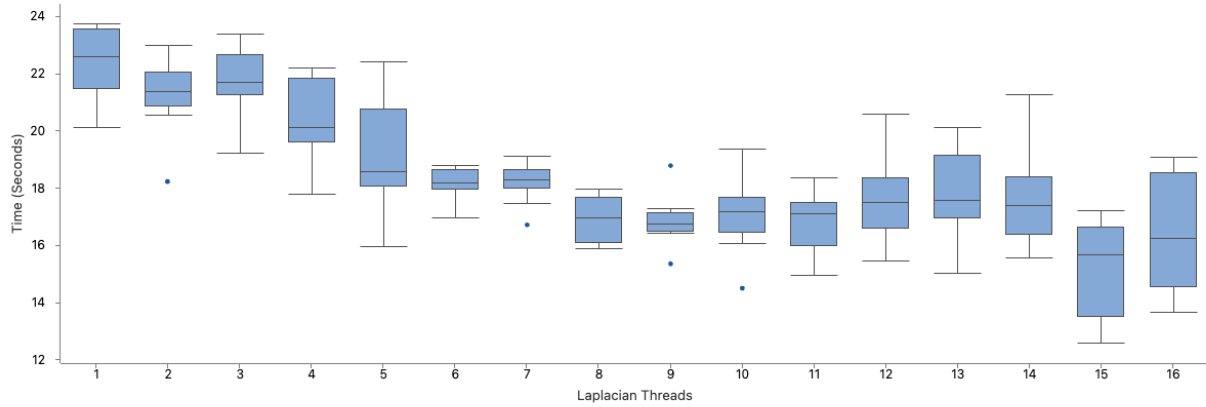

Figure 2: Boxplot of the data.

The boxplot shows the spread in the data from each test performed. Note that from around threads 6 to 9, the data is less spread out, ie has a smaller standard deviation than other values. The dots indicate outliers, these are single data points that lie outside the interquartile range of the data for a single test.
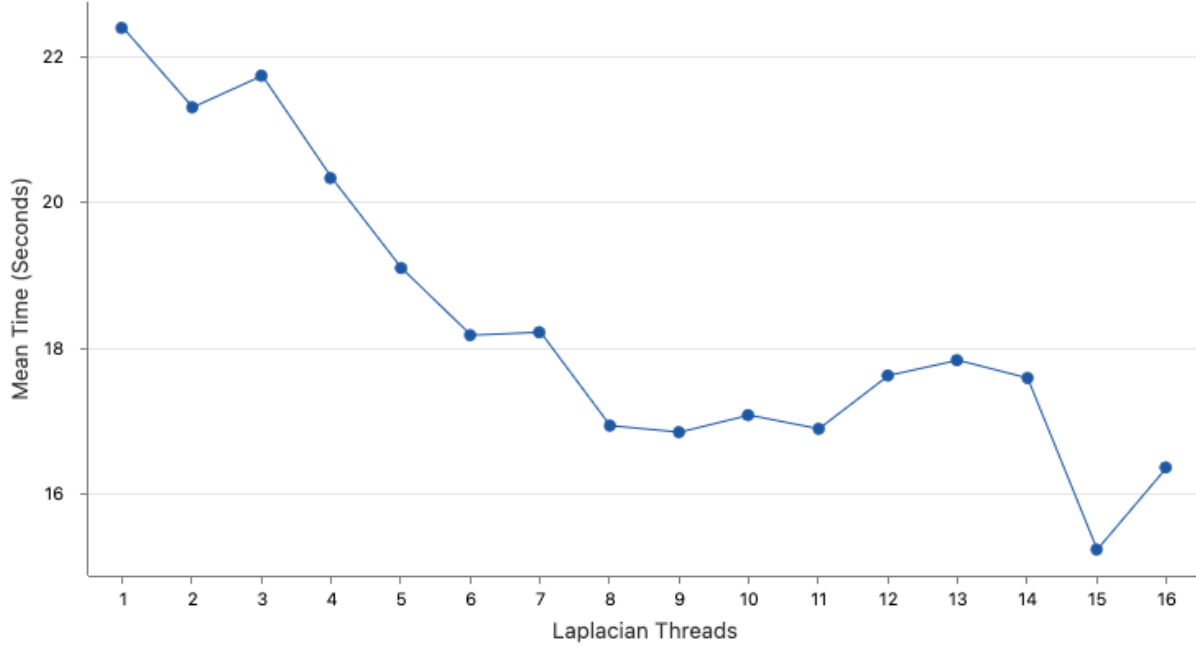
3

Figure 3: Line plot of means.

This lineplot shows the mean for each replicant of each test connected by a line. That is, each dot represents the average of all 10 invocations of `edge_detector` at a given thread amount, from 1 to 16.

Implications of Results

From the boxplot we note that the spread of the data, ie, the standard deviation is lower for the thread counts 6,7,8,9. This could be related to the fact that the processor the experiment was ran on had 6 total cores. However, more experiments on systems with different core counts would be needed to verify this. Thus we cannot say with certainty that this change in spread is anything but a statistical anomaly, but it is an interesting result nonetheless.

The lineplot seems to show the mean time in seconds decreasing as the laplacian threads increase. This increase is most significant from 1 to 6 threads, however, after that is appears to reach an asymptote at around 17 seconds. This seems to indicate that increasing the amount of threads decreases the mean time in seconds up to a certain amount.

Conclusion

Appendix